Project Koios: An audit-first, LLM-assisted workflow for verifiable computational physics

Eugene Joseph M. Ragasa*

Draft article-class preprint — not journal-formatted, peer reviewed, submitted, accepted, or released

Draft status. This article-shaped source reports bounded technical methodology and project experience. Technical review does not constitute peer review, scientific acceptance, release authorization, or publication acceptance.

Abstract

Large-language-model assistance can accelerate the production of plausible code, proofs, numerical analyses, and prose, but plausibility is not scientific authority. We describe an audit-first workflow used across the evolving Project Koios repository family. The workflow freezes a bounded claim, assigns an evidence class, resolves repository and human ownership, binds the claim to immutable public artifacts, applies a checker appropriate to the claim, and records both the supported conclusion and its nonclaims. Five cases illustrate the method: a Lean formalization and exact analytic certificate; provisional optimization/numerics source that defines prospective numerical checks; computational-physics components and provenance in PhysKit; retained finite numerical experiments in `ksdft2effmass`; and candidate Pi/LLM coordination-harness records. The cases preserve failures, corrections, deprecations, and authority boundaries rather than presenting only terminal success. The evidence establishes selected formal, exact, source-level, and bounded numerical facts at identified revisions. It does not establish material validation, universal LLM compliance, productivity benefit, peer review, or publication acceptance.

1 Introduction

Computational-physics work often combines mathematical statements, software, floating-point calculations, scientific interpretation, and decisions about execution or publication. An LLM can contribute to each activity as a drafting or search instrument. The danger is not limited to an incorrect generated answer. A subtler failure occurs when a result that is adequate in one channel is silently promoted into another: a passing test becomes "validated physics," a kernel-checked theorem becomes source fidelity, a repository branch becomes a release, or polished prose becomes publication authority.

The proposed Project Koios architecture separates an agent-facing control plane, a versioned execution plane, and evidence and human-review gates [2]. The reviewed methodology report and its machine-readable evidence ledger freeze the narrower claims used here at public mothership

*Provisional author line. Affiliation, correspondence address, contributor roles, and final authorship metadata are pending explicit human confirmation.

commit `45da16e6428c3c431359da4f80f46de7ebf71081` [5, 8]. This manuscript is an adaptation of that reviewed source, not a replacement for the ledger.

The contribution is methodological and descriptive. We present:

1. a claim-admission workflow in which evidence class is selected before a result is narrated as support;
2. a taxonomy separating formal proof, exact certificates, numerical verification, scientific validation, documented experience, and human decisions;
3. five bounded cases with immutable owner-repository references; and
4. a treatment of failures and corrections as evidence with explicit limits.

The article makes no claim that the complete platform is implemented, that the cases form a representative sample, or that LLM assistance improved time, cost, quality, or productivity.

2 Audit-first method

The workflow begins with the authority boundary rather than the model's capability. Figure 1 summarizes six steps.

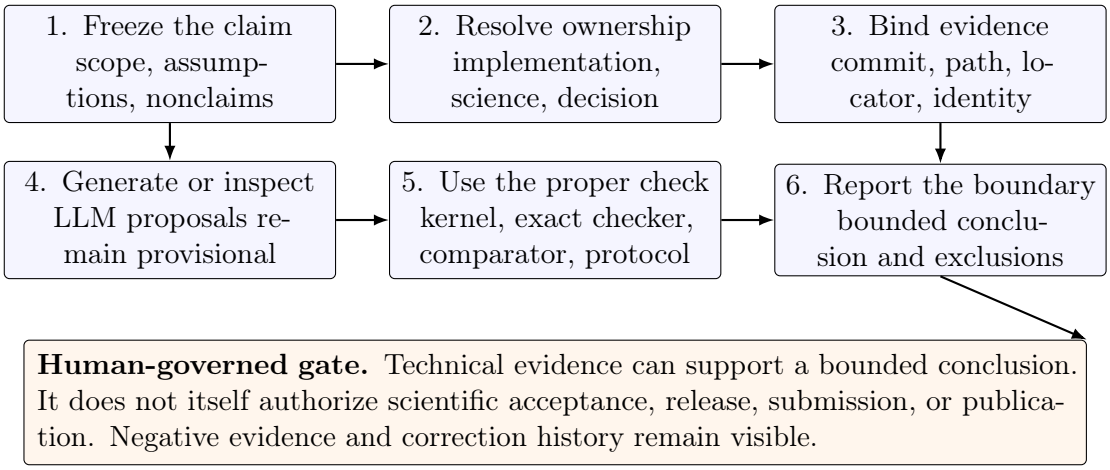


Figure 1: Audit-first claim admission. LLM-produced material remains provisional until the claim is bound to the appropriate evidence and its boundary is reported.

Freeze the claim. The intended conclusion, assumptions, domain, and exclusions are stated before the artifact is treated as evidence. This prevents a successful computation from expanding the claim after the fact.

Resolve ownership. The workflow distinguishes ownership of reusable implementation, retained calculations, scientific meaning, and protected decisions. Repository transfer does not automatically transfer campaign tolerances, execution authority, material interpretation, or publication authority.

Bind immutable evidence. An admitted claim records an owner-repository semantic identity, public remote, commit, repository-relative path, and line or declaration locator. Claim IDs, evidence IDs, and artifact IDs remain distinct, so one claim can cite several artifacts without flattening their roles.

Generate or inspect. An LLM may search bounded sources, propose code or proof terms, draft tests, suggest classifications, and identify inconsistencies. Its output is a proposal, not the acceptance event.

Apply the proper check. A formal proposition is checked by the pinned kernel; an exact certificate requires exact data, a checker, and a proved bridge; a numerical claim requires a declared comparator, tolerance, finite domain, and environment; scientific validation requires scientific reference evidence and review. Human-governed choices remain human decisions.

Report the boundary. The accepted statement is accompanied by explicit nonclaims. Failed attempts, rejected prototypes, unresolved gaps, and deprecations remain visible. The machine-readable ledger implements this structure for the case claims used by the reviewed report [5].

3 Evidence and authority taxonomy

Table 1 summarizes the evidence classes. The classes answer what kind of conclusion an artifact can support; they do not encode release or publication authority.

Table 1: Evidence classes and their principal boundaries.

Class	Bounded basis	Not implied
Implementation conformance	Mapped behavior, contract, and tests at an identified revision	Scientific adequacy, behavior outside the exercised domain, or release approval
Lean theorem evidence	Exact proposition and proof term checked by the pinned kernel, with assumptions and reported axioms	Source fidelity, runtime correctness, numerical convergence, or physical applicability
Exact checked certificate	Exact data accepted by a checker connected by proof to a target proposition	Correctness of unbound external files or floating-point calculations
Numerical verification	Declared finite computation, comparator, tolerance, and environment	Exact equality, continuum behavior, uncertainty quantification, or material validation
Scientific validation	Reference observation or artifact, applicability domain, uncertainty treatment, protocol, and scientific review	Universal validity or publication acceptance
Experience observation	Immutable record of a bounded occurrence and its limitations	Recurrence, general effectiveness, or promotion into another class
Planned work	Explicit future protocol or requirement	Implementation, execution, verification, or completion

Authority is a separate dimension. Claim acceptance, scientific interpretation, execution authorization, release, and publication are not evidence classes. Figure 2 illustrates the separation. A technically adequate artifact can coexist with a pending human decision, and a human decision cannot create a mathematical or numerical result merely by accepting it.

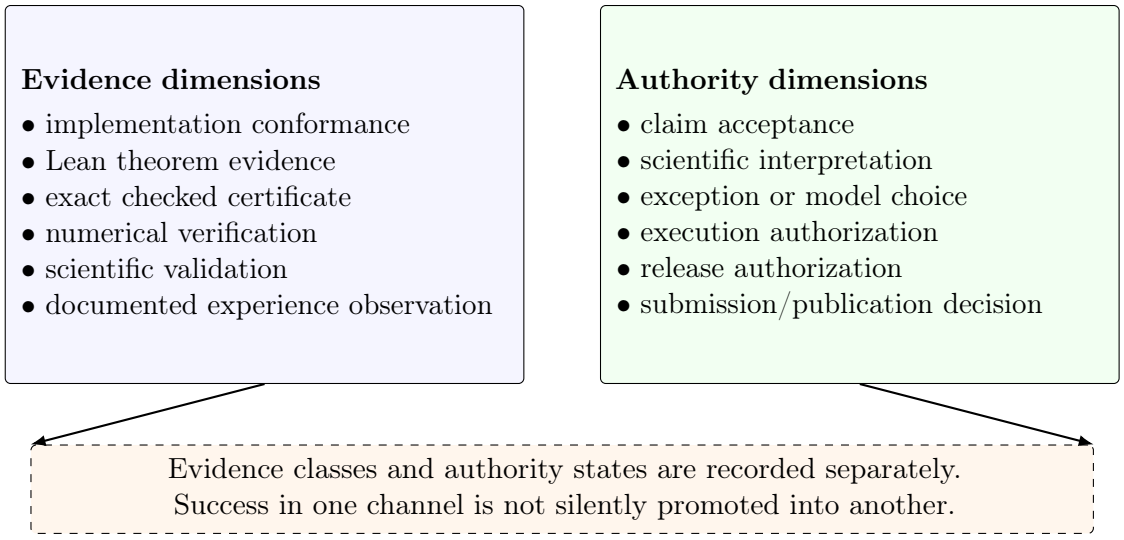


Figure 2: Evidence and authority are recorded as separate dimensions. No arrow denotes automatic promotion from one item to another.

4 Role and limits of LLM assistance

Within the workflow an LLM can assist with discovery, decomposition, drafting, implementation, formal proof search, test construction, and editorial review. The `ksdft2effmass` monograph source discloses assistance with drafting, code, analysis, and editorial work while assigning scientific selection, interpretation, reporting, and responsibility to the author [1]. The cited Formal proof dossier similarly labels a bounded result AI-assisted and kernel-checked [4].

The checker, retained artifact, and accountable person—not the model's rhetorical confidence—carry authority. In particular, an LLM cannot merely by generating text:

- declare its own theorem correct;
- convert approximate output into an exact certificate;
- choose an undisclosed tolerance, model, or applicability domain;
- promote a software test into scientific validation;
- authorize execution, release, submission, or publication; or
- erase a failed attempt and retain only the corrected endpoint.

This article does not expose private transcripts or hidden reasoning. Such material is neither necessary nor sufficient for reproducing the admitted claims. Reproducibility instead rests on public artifacts, immutable identities, and explicit check boundaries. No controlled comparison supports a productivity or quality-benefit claim, and no evidence establishes universal compliance by an LLM with prompts, policy, or scientific norms.

5 Case studies

5.1 Formal: Lean theorem evidence and an exact analytic certificate

At public Formal checkpoint `26603110e9d64cd51c9101eff68060307f258a18`, Lean is pinned to version 4.35.0 and Mathlib to an immutable revision [4]. The encoded two-observation Paper A construction has a proved declaration whose conjunction states zero canonical spectral loss at the encoded witness, membership in the spectral pass set, positive scale, and disjointness from an

optimized-operator pass set. The owner records report the axioms and explicit boundaries.

A separate but connected evidence channel contains exact rational data, a Boolean rational-enclosure checker, checker-soundness theorems, and source-specific bridge theorems. This supports one exact analytic source-certificate chain for the encoded construction. It is important not to collapse the two channels: a Lean theorem and an exact checked certificate are related here, but they answer different audit questions.

Boundary: The checkpoint does not contain a runtime data/parser certificate binding. The exact source certificate does not bind external floating-point artifacts, prove source fidelity, validate a material, or establish a continuum claim. This manuscript does not claim that Lean was rerun during its preparation.

5.2 Optimization and numerics: provisional source definitions

The public migration-branch checkpoint `b722edab216ef416b8b4affa902b04a26e3da98b` documents bounded finite-periodic transforms, evaluation, diagnostics, and optimization capabilities while assigning scientific coordinates, units, thresholds, and acceptance to consumers [6]. Its package metadata records development version `0.1.0.dev0`. It is treated here as provisional source, not a tag, release, or adopted downstream baseline.

Source inspection shows that one finite-series test defines one-, two-, and three-dimensional comparisons against a direct Fourier sum evaluated inside `mpmath.workdps(100)`. Each reference real and imaginary component is converted to binary64 before comparison, and the source defines an explicit floating-point acceptance bound. Other source tests define reconstruction, Parseval, and Hermiticity checks with explicit bounds. A correction commit adds an independently declared Hermiticity-defect case.

These are source-level test definitions for prospective numerical verification. No retained command or CI artifact is admitted as an executed numerical-verification result for this checkpoint. The definitions are not exact certificates or scientific validation, and the library does not select consumer scientific tolerances.

5.3 PhysKit: consistency procedures and provenance boundaries

At public PhysKit checkpoint `11e58938ac8d0022f9e6df3aba7ca8757b4a9f88`, package metadata records development version `0.1.3.dev0`; the checkpoint is not presented as a release [7]. A Morse-potential test source checks force directions and defines a consistency procedure for $F_r = -\mathrm{d}V/\mathrm{d}r$. It compares the returned radial force with a separately expressed centered finite-difference approximation formed from the returned production potential energies, using step 10^{-5} angstrom, relative tolerance 5×10^{-10}, and zero absolute tolerance.

The procedure is not producer-independent: force and energy are returned by the same production evaluator. Moreover, source inspection is not a retained executed result. A periodic-cosine provenance record classifies its mapped checks as implementation and numerical verification while explicitly excluding experimental validation of a material potential.

PhysKit also records a rejected oblique-periodic notebook prototype. Clean execution stopped at a syntax error, saved displays could not be reproduced from the source, and invalid behavior was not migrated. Exact bytes remained recoverable after maintained replacement. This documents review, rejection, correction, and byte-level preservation; it does not document an exercised rollback or a general rollback mechanism.

Extraction records preserve donor identities while retaining campaign tolerances, retained calculations, execution authority, semiconductor-specific interpretation, scientific-validation conclusions,

and human acceptance with `ksdft2effmass`. Reusable represented mathematics can move without silently moving those authorities.

5.4 KSD: retained finite numerical experiments

At public `ksdft2effmass` checkpoint `5aaff25b784c7d1bc079a3030a73be387dc11a0f`, the retained particle-in-a-box case is explicitly an illustrative numerical experiment, not a public API, software harness, semiconductor calculation, or material- validation result [1]. A 38-entry SHA-256 manifest supports maintained-file identity. The reviewed claim uses that manifest only for identity and makes no same-checkpoint producer-script causality claim.

Retained independent verifiers compare finite results against declared finite-difference oracles, projector identities, and residual identities under binary64 arithmetic. Fixed low modes display the expected second-order trend for the declared grids, while modes whose index scales with grid dimension do not show uniform eigenvalue accuracy. This is bounded numerical verification of the finite illustrative construction. It is not evidence about silicon, DFT, Wannier localization, impurity physics, or a material model.

A deprecation record freezes historical Stage C interfaces and exact bytes, names newer adapters, and transfers no scientific, numerical-policy, or execution authority. A separate composite-band successor remains marked `MIGRATION_IN_PROGRESS`; it has no admitted immutable experiment version, maintained run, or new result.

5.5 Pi coordination harness: candidate mechanics and a retained failure

At public Bootstrap checkpoint `7fa09bc3b631984d3b049b00d858935575d5687d`, the accepted incubation boundary permits bounded source, sanitized fixtures, tests, and limitations while excluding transcripts, generated handoffs, credentials, and runtime state [3]. Candidate documents specify deterministic preflight, explicit repository and task naming, fail-closed scoped messaging checks, bounded failures, and the absence of a durable orchestration database. Mocked or fake-boundary test coverage does not prove a live Pi or Herdr service, model compliance, or an authorization service.

A retained incident records that a pane was closed instead of the client being detached, terminating the Pi process. A read-only continuity diagnostic fails closed but cannot recover a terminated process; reassignment remains an authorized human action. The incident supports a bounded experience observation, not autonomous rollback, recurrence, or end-to-end availability.

6 Failures and corrections as evidence

An audit-first account retains negative evidence because it constrains what can be claimed and makes corrections inspectable. Across the cases, the relevant records include:

- an exact formal counterexample as a primary mathematical result rather than a failed success case;
- replacement of a correlated diagnostic assertion by an independently declared test input;
- rejection of a notebook whose saved displays were not reproducible from clean execution;
- deprecation that preserves historical interfaces and byte identities without rewriting provenance; and
- an operational pane-loss incident that motivated a bounded read-only diagnostic rather than an autonomous recovery claim.

These records do not establish a general rollback system, a correction success rate, or a recovery-time distribution. Their narrower contribution is to keep the path from defect to bounded correction visible. Availability of historical bytes is not the same as an exercised rollback.

7 Threats to validity

Selection and survivorship. The cases were selected from one repository family and were not randomly sampled. Repository-visible failures may omit private, informal, abandoned, or undocumented work. No failure rate or general effectiveness estimate follows.

Independence. Many artifacts were authored, reviewed, or selected by the same project owner. A separately expressed comparator can reduce some correlated errors, but common libraries, assumptions, and authorship remain. Kernel checking is not independent scientific peer review. The PhysKit Morse consistency procedure uses values from the same production evaluator and is explicitly not called producer-independent verification.

Evidence freshness and branch semantics. Commits are immutable, but statements about current project state can become stale. Branch reachability does not establish release or adoption. The Optimization case is explicitly provisional, and the KSD commit being reachable from public main does not establish publication or release status.

Tool and model reproducibility. Model behavior can depend on provider, version, prompt, context, tools, and sampling. The method does not promise byte-reproducible model reasoning. It instead freezes accepted public artifacts. Tool availability and dependency resolution can still affect a later rerun.

Scientific reach. The cases concern formal statements, source-level test definitions, finite illustrative calculations, and selected engineering records. They do not validate semiconductor materials, DFT approximations, experimental observables, or multiscale predictive use.

Productivity. No predeclared baseline task set, blinded comparison, complete time/cost record, or controlled study exists. Commit counts, passing tests, token counts, and anecdotal speed cannot support a productivity conclusion.

8 Reproducibility and evidence availability

The reviewed narrative and ledger are publicly addressable at exact Project Koios commit `45da16e6428c3c431359da4f80f46de7ebf71081` [5, 8]. The ledger separates 37 claim identities from 33 evidence records and 37 artifact identities. Each artifact record includes an owner repository, public remote, commit, repository-relative path, and exact line locators. The manuscript cites the owner checkpoints rather than copying their source.

The canonical manuscript source is `Ragasa2026_pk_a.tex`. TikZ figures are retained as reviewable vector source. The local build writes auxiliary files under `.build/`; the rendered PDF retained beside the source is a Draft artifact, not a release. Build-tool identity and the PDF digest are recorded in the publication README.

Reproducing a bounded claim requires consulting the corresponding ledger record, not merely rebuilding this PDF. The PDF is a readable projection of the method; it is not the evidence database and does not promote planned or source-only items into completed results.

9 Rights, authorship, and ethics

The author line is provisional. Affiliation, correspondence, contributor roles, and final authorship require human confirmation. LLM assistance receives neither authorship nor scientific authority merely because text or code was generated. The human author remains responsible for claim selection, interpretation, citations, rights, disclosure, and any future submission decision.

Rights review for a future deposit remains pending. The cited owner repositories contain license and notice records, including repository-specific third-party and adaptation notices. Those records do not establish redistribution permission for every cited or depositable file. This source links to public artifacts and does not vendor owner code, retained scientific data, or protected publication expression.

No private filesystem paths, credentials, session transcripts, hidden runtime state, private benchmark payloads, or unpublished scientific results are included. The work reports software and research-method experience; it does not report a human-subject experiment or use an LLM output as a human acceptance decision. Venue-specific authorship, disclosure, ethics, and rights requirements must be rechecked before any submission.

10 Conclusion

The Project Koios experience supports a narrow methodological conclusion: selected public artifacts demonstrate a practice of separating LLM assistance from formal checking, exact certificate checking, source-level numerical test definitions, retained numerical verification, scientific interpretation, and human-governed decisions. The repositories also preserve selected failures, corrections, and deprecated paths.

The evidence does not establish material validation, universal LLM compliance, productivity benefit, peer review, release, or publication acceptance. The value of the workflow lies in making these boundaries inspectable: each admitted claim is narrower than the set of plausible statements an LLM could draft, and each unresolved authority remains explicit.

References

[1] Eugene Joseph M. Ragasa. ksdft2effmass: retained computational-physics evidence. `https://github.com/eragasa/ksdft2effmass/tree/5aaff25b784c7d1bc079a3030a73be387dc11a0f`, 2026. Immutable public checkpoint 5aaff25b784c7d1bc079a3030a73be387dc11a0f; retained illustrative numerical experiment and provenance records.

[2] Eugene Joseph M. Ragasa. Project koios scientific-platform thesis. `https://github.com/eragasa/projectkoios/blob/965181d89d607fb00469560c94fd50b7a28eaa9d/docs/architecture.scientific-platform-thesis.md`, 2026. Proposed architecture document at immutable commit 965181d89d607fb00469560c94fd50b7a28eaa9d.

[3] Eugene Joseph M. Ragasa. Project koios bootstrap: coordination-harness incubation records. `https://github.com/eragasa/projectkoios-bootstrap/tree/7fa09bc3b631984d3b049b00d858935575d5687d`, 2026. Immutable public checkpoint

7fa09bc3b631984d3b049b00d858935575d5687d; candidate mechanics are not validated generic tooling.

[4] Eugene Joseph M. Ragasa. Project koios formal: certificate vocabulary and theorem evidence. https://github.com/eragasa/projectkoios-formal/tree/26603110e9d64cd51c9101eff68060307f258a18, 2026. Immutable public checkpoint 26603110e9d64cd51c9101eff68060307f258a18; Lean v4.35.0 and pinned Mathlib revision.

[5] Eugene Joseph M. Ragasa. Project koios audit-first methodology evidence ledger. https://github.com/eragasa/projectkoios/blob/45da16e6428c3c431359da4f80f46de7ebf71081/docs/reports/project-koios-audit-first-methodology-evidence.json, 2026. Bounded machine-readable evidence ledger; 37 claims, 33 evidence records, and 37 artifact records.

[6] Eugene Joseph M. Ragasa. Project koios optimization and numerics. https://github.com/eragasa/projectkoios-optimization/tree/b722edab216ef416b8b4affa902b04a26e3da98b, 2026. Public provisional migration-branch checkpoint b722edab216ef416b8b4affa902b04a26e3da98b; development version 0.1.0.dev0, not a release.

[7] Eugene Joseph M. Ragasa. Project koios physkit. https://github.com/eragasa/projectkoios-physkit/tree/11e58938ac8d0022f9e6df3aba7ca8757b4a9f88, 2026. Immutable public checkpoint 11e58938ac8d0022f9e6df3aba7ca8757b4a9f88; development version 0.1.3.dev0, not a release.

[8] Eugene Joseph M. Ragasa. Project koios: An audit-first, llm-assisted workflow for verifiable computational physics. https://github.com/eragasa/projectkoios/blob/45da16e6428c3c431359da4f80f46de7ebf71081/docs/reports/project-koios-audit-first-methodology.md, 2026. Reviewed Draft source at Project Koios commit 45da16e6428c3c431359da4f80f46de7ebf71081; not peer reviewed or released.